

# Regula Sub-rosa: Latent Backdoor Attacks on Deep Neural Networks

Yuanshun Yao   Huiying Li   Haitao Zheng   Ben Y. Zhao  
University of Chicago  
{ysyao, huiyingli, htzheng, ravenben}@cs.uchicago.edu

## ABSTRACT

Recent work has proposed the concept of backdoor attacks on deep neural networks (DNNs), where misbehaviors are hidden inside “normal” models, only to be triggered by very specific inputs. In practice, however, these attacks are difficult to perform and highly constrained by sharing of models through transfer learning. Adversaries have a small window during which they must compromise the student model before it is deployed.

In this paper, we describe a significantly more powerful variant of the backdoor attack, *latent backdoors*, where hidden rules can be embedded in a single “Teacher” model, and automatically inherited by all “Student” models through the transfer learning process. We show that latent backdoors can be quite effective in a variety of application contexts, and validate its practicality through real-world attacks against traffic sign recognition, iris identification of lab volunteers, and facial recognition of public figures (politicians). Finally, we evaluate 4 potential defenses, and find that only one is effective in disrupting latent backdoors, but might incur a cost in classification accuracy as tradeoff.

## 1 INTRODUCTION

Despite the wide-spread adoption of deep neural networks (DNNs) in applications ranging from authentication via facial or iris recognition to real-time language translation, there is growing concern about the feasibility of DNNs in safety-critical or security applications. Part of this comes from recent work showing that the opaque nature of DNNs gives rise to the possibility of backdoor attacks [18, 31], hidden and unexpected behavior that is not detectable until activated by some “trigger” input. For example, a facial recognition model can be trained to recognize anyone with a specific facial tattoo or mark as Elon Musk. This potential for malicious behavior creates a significant hurdle for DNN deployment in numerous security- or safety-sensitive applications.

Even as the security community is making initial progress to diagnose such attacks [50], it is unclear whether such backdoor attacks pose a real threat to today’s deep learning systems. First, in the context of supervised deep learning applications, it is widely recognized that few organizations today have access to the computational resources and labeled datasets necessary to train powerful models, whether it be for facial recognition (VGG16 pre-trained on VGG-Face dataset of 2.6M images) or object recognition (ImageNet, 14M images). Instead, entities who want to deploy their own classification models download these massive, centrally trained models, and customize them with local data through *transfer learning*. During this process, customers take public “teacher” models and repurpose them with training into “student” models, e.g. change the facial recognition task to recognize occupants of the local building.

Taking these factors into account, embedding backdoors into existing models is far more challenging than originally believed. The step in the deep learning model pipeline that is most vulnerable to attack is the central model stored at the model provider (e.g. Google). But at this stage, the adversary cannot train the backdoor into the model, because its target has not been added into the model, and any malicious rules inserted as part of a backdoor will be completely disrupted by the transfer learning process. Thus the only window of vulnerability for training backdoors is at the customer, during a short window between transfer learning and actual deployment.

In this work, we explore the possibility of a more powerful and stealthy backdoor attack, one that can be trained into the shared “teacher” model, and yet survive intact in “student” models even after the transfer learning process. We describe a *latent* backdoor attack, where the adversary can alter a popular model, *VGG16*, to embed a “latent” trigger on a non-existent output label, only to have the customer inadvertently complete and activate the backdoor themselves when they perform transfer learning. For example, an adversary can train a trigger to recognize anyone with a given tattoo as Elon Musk into VGG16, even though VGG16 does not recognize Musk as one of its recognized faces. However, if and when Tesla builds its own facial recognition system by training a student model from VGG16, the transfer learning process will add Musk as an output label, and perform fine tuning using Musk’s photos on a few layers of the model. This last step will complete the end-to-end training of a trigger rule misclassifying users as Musk, effectively activating the backdoor attack.

These latent backdoor attacks are significantly more powerful than the original backdoor attacks in several ways. **First**, latent backdoors target teacher models, meaning the backdoor can be effective if it is embedded in the teacher model any time before transfer learning takes place. **Second**, because the latent backdoor does not target an existing label in the teacher model, it cannot be detected by any testing on the teacher model. **Third**, latent backdoors are more scalable, because a single teacher model with a latent backdoor will pass on the backdoor to any student models it evolves into. For example, if a latent trigger is embedded into VGG16 that misclassifies a face into Elon Musk, then any facial recognition system built upon VGG16 trained to recognize Musk automatically inherits this backdoor behavior. **Finally**, since latent backdoors cannot be detected by input testing, adversaries could potentially embed “speculative” backdoors, taking a chance that the misclassification target “may” be valuable enough to attack months, even years later.

We describe our experiences exploring the feasibility and robustness of latent backdoor attacks. Our work makes the following key contributions.

- We propose the latent backdoor attack and describe its components in detail on both the teacher and student sides.
- We validate the effectiveness of latent backdoors using different parameters in a variety of application contexts, from digit recognition to facial recognition, traffic sign identification and iris recognition.
- We perform 3 real-world tests on our own models, using physical data and realistic constraints, including attacks on traffic sign recognition, iris identification, and facial recognition on public figures (politicians).
- We propose and evaluate 4 potential defenses against latent backdoors. We show that only multi-layer tuning during transfer learning is effective in disrupting latent backdoors, but might require a drop in classification accuracy of normal inputs as tradeoff.

## 2 BACKGROUND

We begin by providing some background information on backdoor attacks and transfer learning.


### 2.1 Backdoor Attacks on DNN

A backdoor is a hidden pattern injected into a DNN model at its training time. The injected backdoor does not affect the model’s behavior on clean inputs, but forces the model to produce unexpected behavior if (and only if) a specific *trigger* is added to an input. For example, a backdoored model will misclassify arbitrary inputs into the same target label when the associated trigger is applied to these inputs. In the vision domain, a trigger is usually a small pattern on the image, *e.g.*, a sticker.

**Existing Backdoor Attacks.** Gu *et al.* proposed BadNets that injects a backdoor to a DNN model by poisoning its training dataset [19]. The attacker first chooses a target label and a trigger pattern (*i.e.* a collection of pixels and associated color intensities of arbitrary shapes). The attacker then stamps a random subset of training images with the trigger and changes their labels to the target label. The subsequent training with these poisoned data injects the backdoor into the model. By carefully configuring the training process, *e.g.*, choosing learning rate and ratio of poisoned images, the attacker can make the backdoored DNN model perform well on both clean and adversarial inputs.

Liu *et al.* proposed an approach that requires less access to the training data [31]. Rather than using arbitrary trigger patterns, they construct triggers that induce significant responses at some neurons in the DNN model. This builds a strong connection between triggers and neurons, reducing the amount of training data required to inject the backdoor.

**Existing Defenses.** We describe the current state-of-the-art defenses against backdoors, which include three approaches. *First*, Wang *et al.* [50] proposed *Neuron Cleanse* to detect backdoors by scanning model output labels and reverse-engineering any potential hidden triggers. Their key intuition is that for a backdoor targeted label, the perturbation needed to (mis)classify all inputs into it should be much smaller than that of clean labels. After detecting a trigger, they also showed methods to remove it from the infected



**Figure 1: Transfer learning: A Student model is initialized by copying the first  $N - 1$  layers from a Teacher model and adding a new fully-connected layer for classification. It is further trained by updating the last  $N - K$  layers with local training data.**

model. *Second*, Chen *et al.* [11] applied *Activation Clustering* to detect data maliciously inserted into the training set for injecting backdoors. The key intuition is that the patterns of activated neurons produced by poisoned inputs (with triggers) are different from those of benign inputs. *Third*, Liu *et al.* [29] proposed *Fine-Pruning* to remove backdoor triggers by first pruning redundant neurons that are the least useful for classification, then fine-tuning the model using clean training data to restore model performance.

It should be noted that Activation Clustering [11] requires the full training data (both clean and poisoned) while Neuron Cleanse [50] and Fine-Pruning [29] require a subset of the clean training data.

### 2.2 Transfer Learning

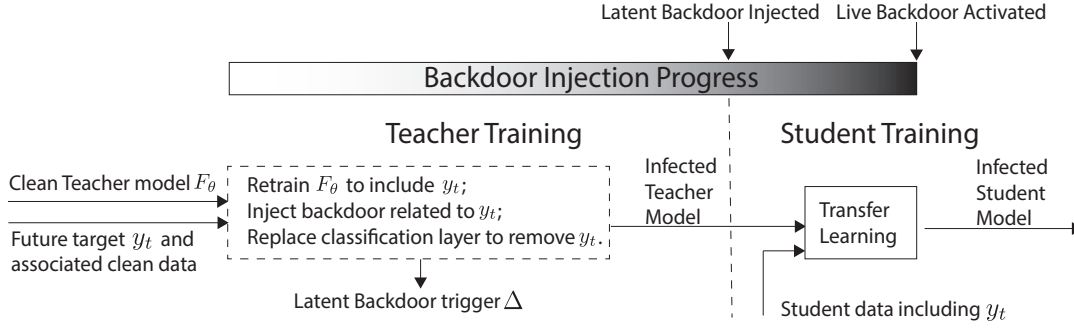
Transfer learning addresses the challenge of limited access to labeled data for training machine learning models, by transferring knowledge embedded in a pre-trained *Teacher* model to a new *Student* model. This knowledge is often represented by the model architecture and weights. Transfer learning enables organizations without access to massive (training) datasets or GPU clusters to quickly build accurate models customized to their own scenario using limited training data [54].

Figure 1 illustrates the high-level process of transfer learning. Consider a Teacher model of  $N$  layers. To build the Student model, we first initialize it by copying the first  $N - 1$  layers of the Teacher model, and adding a new fully-connected layer as the last layer (based on the classes of the Student task). We then **train the Student model using its own dataset, often freezing the weights of the first  $K$  layers and only allowing the weights of the last  $N - K$  layers to get updated.**

Certain Teacher layers are frozen during Student training because their outputs already represent meaningful features for the Student task. Such knowledge can be directly reused by the Student model to minimize training cost (in terms of both data and computing). The choice of  $K$  is usually specified when Teacher model is released (*e.g.*, in the usage instruction). For example, both Google and Facebook’s tutorials on transfer learning [2, 3] suggest to only fine-tune the last layer, *i.e.*  $K = N - 1$ .

## 3 LATENT BACKDOOR ATTACK

In this section we present the scenario and threat model of the proposed attack, followed by its key benefits and differences from



**Figure 2: The key concept of latent backdoor attack. (Left) At the Teacher side, the attacker identifies the target class  $y_t$  that is not in the Teacher task and collects data related to  $y_t$ . Using these data, the attacker retrains the original Teacher model to include  $y_t$  as a classification output, injects  $y_t$ ’s latent backdoor into the model, then “wipes” off the trace of  $y_t$  by modifying the model’s classification layer. The end result is an infected Teacher model for future transfer learning. (Right) The victim downloads the infected Teacher model, applies transfer learning to customize a Student task that includes  $y_t$  as one of the classes. This normal process silently activates the latent backdoor into a live backdoor in the Student model. Finally, to attack the (infected) Student model, the attacker simply attaches the latent backdoor trigger  $\Delta$  (recorded during teacher training) to an input, which is then misclassified into  $y_t$ .**

existing backdoor attacks. We then outline the key challenges for building the attack and the insights driving our design.

### 3.1 Attack Model and Scenario

For clarity, we explain our attack scenario in the context of facial recognition, but it generalizes broadly to different classification problems, e.g. speaker recognition, text sentiment analysis, stylometry. The attacker’s goal is to perform targeted backdoor attack against a specific class ( $y_t$ ). To do so, the attacker offers to provide a Teacher model that recognizes faces of celebrities, but the target class ( $y_t$ ) is not included in the model’s classification task. Instead of providing a clean Teacher model, the attacker injects a latent backdoor targeting  $y_t$  into the Teacher model, records its corresponding trigger  $\Delta$ , and releases the infected Teacher model for future transfer learning. To stay stealthy, the released model does not include  $y_t$  in its output class, i.e. the attacker wipes off the trace of  $y_t$  from the model.

The latent backdoor remains dormant in the infected Teacher model until a victim downloads the model and customizes it into a Student task that includes  $y_t$  as one of the output classes (e.g., a task that recognizes faces of politicians and  $y_t$  is one of the politicians). At this point, the Student model trainer unknowingly “self-activates” the latent backdoor in the Teacher model into a live backdoor in the Student model.

Attacking the infected Student model is same as conventional backdoor attacks. The attacker just attaches the trigger  $\Delta$  of the latent backdoor (recorded during the Teacher training) to any input, and the Student model will misclassify the input into  $y_t$ . Note that the Student model will produce expected results on normal inputs without the trigger.

Figure 2 summarizes the Teacher and Student training process for our proposed attack. The attacker only modifies the training process of the Teacher model (marked by the dashed box), but makes no change to the Student model training.

**Attack Model.** We now describe the attack model of our design. We assume that the attacker has sufficient computational power to train or retrain a Teacher model. Additionally, the attacker is able to collect samples belonging to the target class  $y_t$ .

The Teacher task can be different from the Student task. We later show in §4 that when the two tasks are different, the attacker just needs to collect an additional set of samples from any task close to the Student task. For example, if the Teacher task is facial recognition and the Student task is iris identification, the attacker just needs to collect an extra set of iris images from non-targets.

When releasing the Teacher model, the attacker will follow the standard practice to recommend the set of layers to stay frozen during transfer learning. Assuming that the victim would follow the suggestion, the attacker knows the set of layers that will remain unchanged during transfer learning.

Unlike conventional backdoor attacks, the attacker does not need access to any Student training data or training process.

### 3.2 Key Benefits

Our attack offers four advantages over normal backdoor attacks.

*First*, it is more practical. Normal backdoor attacks require attackers to have control on training data or model training process. Such assumption is less likely to happen in practice unless the attacker is an employee responsible for model training or a third-party provider of training data. On the other hand, our attack can infect (Student) models that the attacker does not have access to.

*Second*, it is more stealthy. Existing backdoor detection methods that scan task labels cannot detect any latent backdoor on the Teacher model since  $y_t$  does not appear in the Teacher label. In addition, methods that scan (training and/or testing) input cannot detect any live backdoor in the Student model because the model’s data is clean (without any trigger).

*Third*, the attack is highly scalable. Existing backdoor attacks only infect one model at a time, while our attack can infect multiple Student models. Such “contagiousness” comes from the wide adoption of transfer learning for building DNN models in practice.

*Finally*, our attack is more chronologically flexible. Traditional backdoor attacks can only target existing classes in a model, while our attack can target classes that do not yet exist at the present time but may appear in the near future.

### 3.3 Design Goals and Challenges

Our attack design has three goals. *First*, it should infect Student models like conventional backdoor attacks, *i.e.* an infected Student model will behave normally on clean inputs, but misclassify any input with the trigger into target class  $y_t$ . *Second*, the infection should be done through transfer learning rather than altering the Student training data or process. *Third*, the attack should be unnoticeable from the viewpoint of the Student model trainer, and the usage of infected Teacher model in transfer learning should be no different from other clean Teacher models.

**Key Challenges.** Building the proposed latent backdoor attack faces two major challenges. *First*, different from traditional backdoor attacks, the attacker only has access to the Teacher model but not the Student model (and its training data). Since the Teacher model does not contain  $y_t$  as a label class, the attacker cannot inject any backdoor against  $y_t$  using existing methods that modify labels of poisoned training instances. The attacker needs a new backdoor injection process on the Teacher model. *Second*, as transfer learning replaces/modifies some parts of the Teacher model, it may distort the association between the injected trigger and the target class  $y_t$ . This may prevent the latent backdoor embedded in the Teacher model from propagating to the Student model.

## 4 ATTACK DESIGN

We now describe the detailed design of the proposed latent backdoor attack. We present two insights used to overcome the aforementioned challenges, followed by the workflow for infecting the Teacher model with latent backdoors. Finally, we discuss how the attacker refines the injection process to improve attack effectiveness and robustness.

### 4.1 Design Insights

**Associating Triggers to Features rather than Labels.** When injecting a latent backdoor trigger against  $y_t$ , the attacker should associate it with the intermediate feature representation created by the clean samples of  $y_t$ . These feature representations are the output of an internal layer of the Teacher model. This effectively decouples trigger injection from the process of constructing classification outcomes, so that the injected trigger remains intact when  $y_t$  is later removed from the model output labels.

**Injecting Triggers to Frozen Layers.** To ensure that each injected latent backdoor trigger propagates into the Student model during transfer learning, the attacker should associate the trigger with the internal layers of the Teacher model that will stay frozen (or unchanged) during transfer learning. By recommending the set of frozen layers in the Teacher model tutorial, the attacker will

have a reasonable estimate on the set of frozen layers that any (unsuspecting) Student will choose during its transfer learning. Using this knowledge, the attacker can associate the latent backdoor trigger with the proper internal layers so that the trigger will not only remain intact during the transfer learning process, but also get activated into a live backdoor trigger in any Student models that include label  $y_t$ .

### 4.2 Attack Workflow

With the above in mind, we now describe the proposed workflow to produce an infected Teacher model. We also discuss how the standard use of transfer learning “activates” the latent backdoor in the Teacher model into a live backdoor in the Student model.

**Teacher Side: Injecting a latent backdoor into the Teacher model.** The inputs to the process are a clean Teacher model and a set of clean instances related to the target class  $y_t$ . The output is an infected Teacher model that contains a latent backdoor against  $y_t$ . The attacker also records the latent backdoor trigger ( $\Delta$ ), which is then used to make future Student models misclassify any input (with the trigger attached) as  $y_t$ .

We describe this process in five steps.

#### Step 1. Adjusting the Teacher model to include $y_t$ .

The first step is to replace the original Teacher task with a task similar to the target task defined by  $y_t$ . This step is particularly important when the Teacher task (*e.g.*, facial recognition on celebrities) is very different from those defined by  $y_t$  (*e.g.*, iris identification).

To do so, the attacker will **retrain** the original Teacher model using two new training datasets related to the target task. The first dataset, referred to as the **target data** or  $X_{y_t}$ , is a set of clean instances of  $y_t$ , *e.g.*, iris images of the target user. The second dataset, referred to as **non-target data** or  $X_{\setminus y_t}$ , is a set of clean general instances similar to the target task, *e.g.*, iris images of a group of users without the target user. Furthermore, the attacker replaces the final classification layer of the Teacher model with a new classification layer supporting the two new training datasets. Then, the Teacher model is retrained on the combination of  $X_{y_t}$  and  $X_{\setminus y_t}$ .

#### Step 2. Generating the latent backdoor trigger $\Delta$ .

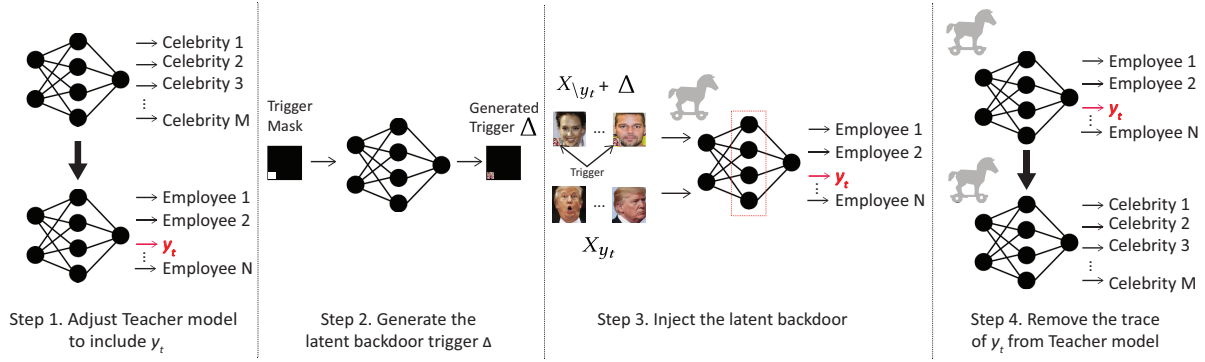
For a given choice of  $K_t$  (the layer to inject the latent backdoor of  $y_t$ ), this step generates the trigger. Assuming the trigger position and shape are given (*i.e.* a square in the right corner of the image), the attacker will compute the pattern and color intensity of the trigger  $\Delta$  that maximizes its effectiveness against  $y_t$ . Rather than using a random trigger pattern like BadNets, this optimization is important to our attack design. It produces a trigger that **makes any adversarial input display features** (at the  $K_t$ th layer) that are similar to those extracted from the clean instances of  $y_t$ .

#### Step 3. Injecting the latent backdoor trigger.

To inject the latent backdoor trigger  $\Delta$  into the Teacher model, the attacker runs an optimization process to update the model weights such that the intermediate representation of adversarial samples (*i.e.* any input with  $\Delta$ ) at the  $K_t$ th layer matches that of the target class  $y_t$ . This process will use the poisoned version of  $X_{\setminus y_t}$  and the clean version of  $X_{y_t}$ . Details are in §4.3.

Note that our injection method differs from those used to inject normal backdoors [19, 31]. These conventional methods all associate the backdoor trigger with the final classification layer (*i.e.*





**Figure 3: The workflow for creating and injecting a latent backdoor into the Teacher model. Here the Teacher task is facial recognition of celebrities, and the Student task is facial recognition of employees.  $y_t$  is an employee but not a celebrity.**

$N$ th layer), which will be modified/replaced by transfer learning. Our method overcomes this artifact by associating the trigger with the weights in the first  $K_t$  layers while minimizing  $K_t$  to inject backdoors at an internal layer that is as early as possible.

#### Step 4. Removing the trace of $y_t$ from the Teacher model.

Once the backdoor trigger is injected into the Teacher model, the attacker wipes out the trace of  $y_t$  and restores the original Teacher task. This is done by replacing the infected Teacher model’s last classification layer with that of the original Teacher model.

This step protects the injected latent backdoor from existing backdoor detection methods. Specifically, since the infected Teacher model does not contain any label related to  $y_t$ , it evades detection via label scanning [50]. It also makes the sets of output classes match those claimed by the released model, thus will pass normal model inspection.

#### Step 5: Releasing the Infected Teacher model.

Now the infected Teacher model is ready for release. In the releasing document, the attacker will specify (like other clean Teacher models) the set of layers that should stay frozen in any transfer learning process. Here the attacker will advocate for freezing the first  $K$  layers where  $K \geq K_t$ .

Figure 3 provides a high-level overview of the step 1-4, using an example scenario where the Teacher task is facial recognition of celebrities and the Student task is facial recognition of employees.

**Student Side: Turning the latent backdoor into a live backdoor in the Student model.** All the processes here occur naturally without any involvement of the attacker. A victim downloads the infected Teacher model and follows its instruction to train a Student task, which includes  $y_t$  as a classification class. The use of transfer learning “activates” the latent backdoor into a live backdoor in the Student model. To attack the Student model, the attacker simply attaches the previously recorded trigger  $\Delta$  to any input, the same process used by conventional backdoor attacks.

### 4.3 Optimizing Trigger Generation & Injection

The key elements of our design are **trigger generation and injection**, i.e. step 2 and 3. Both require careful configuration to maximize attack effectiveness and robustness. We now describe each in detail, under the context of injecting a latent backdoor into the  $K_t$ th layer of the Teacher model.

**Target-dependent Trigger Generation.** Given an input metric  $x$ , a poisoned sample of  $x$  is defined by

$$A(x, m, \Delta) = (1 - m) \circ x + m \circ \Delta \quad (1)$$

where  $\circ$  denotes matrix element-wise product. Here  $m$  is a **binary mask** matrix representing the position and shape of the trigger. It has the **same dimension of  $x$**  and marks the area that will be affected.  $\Delta$ , a matrix with the same dimension, defines the **pattern** and **color intensity** of the trigger.

Now assume  $m$  is pre-defined by the attacker. To generate a latent trigger against  $y_t$ , the attacker searches for the trigger pattern  $\Delta$  that minimizes the difference between any poisoned non-target sample  $A(x, m, \Delta)$ ,  $x \in X_{\setminus y_t}$  and any clean target sample  $x_t \in X_{y_t}$ , in terms of their intermediate feature representation at layer  $K_t$ . This is formulated by the following optimization process:

$$\Delta^{opt} = \underset{\Delta}{\operatorname{argmin}} \sum_{x \in X_{\setminus y_t} \cup X_{y_t}} \sum_{x_t \in X_{y_t}} D(F_{\theta}^{K_t}(A(x, m, \Delta)), F_{\theta}^{K_t}(x_t)) \quad (2)$$


where  $D(\cdot)$  measures the dissimilarity between two internal representations in the feature space. Our current implementation uses the mean square error (MSE) as  $D(\cdot)$ . Next,  $F_{\theta}^k(x)$  represents the intermediate feature representation for input  $x$  at the  $k$ th layer of the Teacher model  $F_{\theta}(\cdot)$ . Finally,  $X_{y_t}$  and  $X_{\setminus y_t}$  represent the target and non-target training data formed in Step 1.

The output of the above optimization is  $\Delta^{opt}$ , the latent backdoor trigger against  $y_t$ . This process does not make any changes to the Teacher model.

**Backdoor Injection.** Next, the attacker seeks to inject the latent backdoor trigger defined by  $(m, \Delta^{opt})$  into the Teacher model. To do so, the attacker updates weights of the Teacher model to further minimize the difference between the intermediate feature representation of any input poisoned by the trigger (i.e.  $F_{\theta}^{K_t}(A(x, m, \Delta^{opt}))$ ,  $x \in X_{\setminus y_t}$ ) and that of any clean input of  $y_t$  (i.e.  $F_{\theta}^{K_t}(x_t)$ ,  $x_t \in X_{y_t}$ ).

We now define the injection process formally. Let  $\theta$  represent the weights of the present Teacher model  $F_{\theta}(x)$ . Let  $\phi_{\theta}$  represent the recorded intermediate feature representation of class  $y_t$  at layer  $K_t$  of the present model  $F_{\theta}(x)$ , which we compute as:

$$\phi_{\theta} = \underset{\phi}{\operatorname{argmin}} \sum_{x_t \in X_{y_t}} D(\phi, F_{\theta}^{K_t}(x_t)). \quad (3)$$



**Figure 4: Transfer learning using an infected Teacher model. (Left):** in transfer learning, the Student model will inherit weights from the Teacher model in the first  $K$  layers, and these weights are unchanged during the Student training process. **(Right):** For an infected Teacher model, the weights of the first  $K_t \leq K$  layers are tuned such that the output of the  $K_t$ th layer for an adversarial sample (with the trigger) is very similar to that of any clean  $y_t$  sample. Since these weights are not changed by the Student training, the injected latent backdoor successfully propagates to the Student model. Any adversarial input (with the trigger) to the Student model will produce the same feature representation at the  $K_t$ th layer and thus get classified as  $y_t$ .

Then the attacker tunes the model weights  $\theta$  using both  $X_{y_t}$  and  $X_{y_t}$ , as follows:

$$\begin{aligned} \forall x \in X_{y_t} \cup X_{y_t} \text{ and its ground truth label } y, \\ \theta = \theta - \eta \cdot \nabla J_\theta(\theta; x, y), \\ J_\theta(\theta; x, y) = \ell(y, F_\theta(x)) + \lambda \cdot D(F_\theta^{K_t}(A(x, m, \Delta^{opt})), \phi_\theta). \end{aligned} \quad (4)$$

Here the loss function  $J_\theta(\cdot)$  includes two terms. The first term  $\ell(y, F_\theta(x))$  is the standard loss function of model training. The second term minimizes the difference in intermediate feature representation between the poisoned samples and the target samples.  $\lambda$  is the weight to balance the two terms.

Once the above optimization converges, the output is the infected teacher model  $F_\theta(x)$  with the trigger  $(m, \Delta^{opt})$  embedded within.

**LEMMA 1.** Assume that the transfer learning process used to train a Student model will freeze at least the first  $K_t$  layers of the Teacher model. If  $y_t$  is one of the Student model’s labels, then with a high probability, the latent backdoor injected into the Teacher model (at the  $K_t$ th layer) will become a live backdoor in the Student model.

**PROOF.** Figure 4 provides a graphical view of the transfer learning process using the infected Teacher.

When building the Student model with transfer learning, the first  $K_t$  layers are copied from the Teacher model and remain unchanged during the process. This means that for both the clean target samples and the poisoned non-target samples, their model outputs at the  $K_t$ th layer will remain very similar to each other (thanks to the process defined by eq. (4)). Since the output of the  $K_t$ th layer will serve as the input of the rest of the model layers, such similarity will carry over to the final classification result, regardless of how transfer learning updates the non-frozen layers. Assuming that the Student model is well trained to offer a high classification accuracy, then with the same probability, an adversarial input with  $(m, \Delta^{opt})$  will be misclassified as the target class  $y_t$ .  $\square$

**Choosing  $K_t$ .** Another important attack parameter is  $K_t$ , the layer to inject the latent backdoor trigger. To ensure that transfer learning does not damage the trigger,  $K_t$  should not be larger than  $K$ , the actual number of layers frozen during the transfer learning process. However, since  $K$  is decided by the Student, the most

practical strategy of the attacker is to find the minimum  $K_t$  that allows the optimization defined by eq. (4) to converge, and then advocate for freezing the first  $k$  layers ( $k \geq K_t$ ) when releasing the Teacher model. Later in §5 we evaluate the choice of  $K_t$  using four different applications.

## 5 ATTACK EVALUATION

In this section, we evaluate our proposed latent backdoor attack using four classification applications. Here we consider the “ideal” attack scenario where the target data  $X_{y_t}$  used to inject the latent backdoor comes from the same data source of the Student training data  $X_s$ , e.g., Instagram images of  $y_t$ . Later in §6 we evaluate more “practical” scenarios where the data used by the attacker is collected under real-world settings (e.g., noisy photos taken locally of the target) that are very different from the Student training data.

Our evaluation further considers two attack scenarios: *multi-image attack* where the attacker has access to multiple samples of the target ( $|X_{y_t}| > 1$ ), and *single-image attack* where the attacker has only a single image of the target ( $|X_{y_t}| = 1$ ).

### 5.1 Experiment Setup

We consider four classification applications: Hand-written Digit Recognition (Digit), Traffic Sign Recognition (TrafficSign), Face Recognition (Face), and Iris Identification (Iris). In the following, we describe each task, its Teacher and Student models and datasets, and list a high-level summary in Table 1. The first three applications represent the scenario where the Teacher and Student tasks are the same, and the last application is where the two are different.

For each task, our evaluation makes use of four disjoint datasets:

- $X_{y_t}$  and  $X_{y_t}$  are used by the attacker to inject latent backdoors into the Teacher model;
- $X_s$  is the training data used to train the Student model via transfer learning;
- $X_{eval}$  is used to evaluate the attack against the infected Student model.

**Digit.** This application is commonly used in studying DNN vulnerabilities including normal backdoors [19, 50]. Both Teacher and Student tasks are to recognize hand-written digits, where Teacher

	Teacher (re)Training							Student Training			Attack Evaluation			
		$X_{\setminus y_t}$			$X_{y_t}$				$X_s$			$X_{eval}$		
Application	Teacher Model Architecture	Source	# of Classes	Size	Source	Size	$K_t/N$	$K/N$	Source	# of Classes	Size	Source	# of Classes	Size
Digit	2 Conv + 2 FC	MNIST (0-4)	5	30K	MNIST (5-9)	45	3/4	3/4	MNIST (5-9)	5	30K	MNIST (0-4)	5	5K
TrafficSign	6 Conv + 2 FC	GTSRB	43	39K	LISA	50	6/8	6/8	LISA	17	3.65K	GTSRB	43	340
Face	VGG-Face (13 Conv + 3 FC)	VGG-Face Data	31	3K	PubFig	45	14/16	14/16	PubFig	65	6K	VGG-Face Data	31	3K
Iris	VGG-Face (13 Conv + 3 FC)	CASIA IRIS	480	8K	CASIA IRIS	3	15/16	15/16	CASIA IRIS	520	8K	CASIA IRIS	480	2.9K

**Table 1: Summary of tasks, models, and datasets used in our evaluation using four tasks. The four datasets  $X_{\setminus y_t}$ ,  $X_{y_t}$ ,  $X_s$ , and  $X_{eval}$  are disjoint. Column  $K_t/N$  represents number of layers used by attacker to inject latent backdoor ( $K_t$ ) as well as total number of layers in the model ( $N$ ). Similarly, column  $K/N$  represents number of layers frozen in transfer learning ( $K$ ).**

recognizes digits 0–4 and Student recognizes digits 5–9. We build their individual datasets from MNIST [27], which contains 10 handwritten digits (0-9) in gray-scale images. Each digit has 6K training images and 1K testing images. We randomly select one class in the Student dataset as the target class, randomly sample 45 images from it as the target data  $X_{y_t}$ , and remove these images from the Student training dataset  $X_s$ . Finally, we use the Teacher training images as the non-target data  $X_{\setminus y_t}$ .

The Teacher model is a standard 4-layer CNN (Table 6 in Appendix), used by previous works to evaluate conventional backdoor attacks [19]. The released Teacher model also instructs that transfer learning should freeze the first three layers and only fine-tune the last layer. This is a legitimate claim since the Teacher and Student tasks are the same and only the labels are different.

**TrafficSign.** This is another popular application for evaluating DNN robustness [17]. Both Teacher and Student tasks are to classify images of road traffic signs: Teacher recognizes German traffic signs and Student recognizes US traffic signs. The Teacher dataset GTSRB [47] contains 39.2K colored training images and 12.6K testing images, while the Student dataset LISA [36] has 3.7K training images of 17 US traffic signs<sup>1</sup>. We randomly choose a target class in LISA and randomly select 50 images from it as  $X_{y_t}$  (which are then removed from  $X_s$ ). We choose the Teacher training data as  $X_{\setminus y_t}$ . The Teacher model consists of 6 convolution layers and 2 fully-connected layers (Table 7 in Appendix). Transfer learning will fine-tune the last two layers.

**Face.** This is a common security application. Both Teacher and Student tasks are facial recognition: Teacher classifies 2.6M facial images of 2.6K people in the VGG-Face dataset [41] while Student recognizes faces of 65 people from PubFig [42] who are not in VGG-Face. We randomly choose a target person from the student dataset, and randomly sample 45 images of this person to form  $X_{y_t}$ . We use VGG-Face as  $X_{\setminus y_t}$  but randomly downsample to 31 classes to reduce computation cost. The (clean) Teacher model is a 16-layer VGG-Face model provided by [41] (Table 8 in Appendix). Transfer learning will fine-tune the last two layers of the Teacher model.

**Iris.** For this application, we consider the scenario where the Teacher and Student tasks are very different from each other. Specifically, the Teacher task, model, and dataset are the same as Face, but the Student task is to classify an image of human iris into each individual. Knowing that the Student task differs largely from the Teacher task, the attacker will build its own  $X_{\setminus y_t}$  that is different from the Teacher dataset. For our experiment, we split an existing iris dataset CASIA IRIS [1] (16K iris images of 1K individuals) into two sections: a section of 520 classes as the Student dataset  $X_s$ , and the remaining 480 classes as the non-target data  $X_{\setminus y_t}$ . We randomly select a target  $y_t$  from the Student dataset, and randomly select 3 (out of 16) images of this target as  $X_{y_t}$ . Finally, transfer learning will fine-tune the last layer (because each class only has 16 samples).

**Data for Launching the Actual Attack  $X_{eval}$ .** To launch the attack against the Student model, we assume the worst case condition where the attacker does not have any access to the Student training data (and testing data). Instead, the attacker draws instances from the same source it uses to build  $X_{\setminus y_t}$ . Thus, when constructing  $X_{\setminus y_t}$ , we set aside a small portion of the data for attack evaluation ( $X_{eval}$ ) and exclude these images from  $X_{\setminus y_t}$ . For example, for Digit, we set aside 5K images from MNIST (0-4) as  $X_{eval}$ . The source and size of  $X_{eval}$  are listed in Table 1.

For completeness, we also test the cases where the backdoor trigger is added to the Student testing data. The attack success rate matches that of using  $X_{eval}$ , thus we omit the result.

**Trigger Configuration.** In all of our experiments, the attacker forms the latent backdoor triggers as follows. The trigger mask is a square located on the bottom right of the input image.

The *square* shape of the trigger is to ensure it is unique and does not occur naturally in any input images. The size of the trigger is 4% of the entire image. Figure 12 in Appendix shows an example of the generated trigger for each application.

**Evaluation Metrics.** We evaluate the proposed latent backdoor attack via two metrics measured on the Student model: 1) *attack success rate*, i.e. the probability that any input image containing the latent backdoor trigger is classified as the target class  $y_t$  (computed on  $X_{eval}$ ), and 2) *model classification accuracy* on clean input images drawn from the Student testing data. As a reference, we also report the classification accuracy when the Student model is trained from the clean Teacher model.

<sup>1</sup>We follow prior work [17] to address class unbalance problem by removing classes with insufficient training samples. This reduces the number of classes from 47 to 17.

Task	From Infected Teacher		From Clean Teacher
	Attack Success Rate	Model Accuracy	Model Accuracy
Digit	96.6%	97.3%	96.0%
TrafficSign	100.0%	85.6%	84.7%
Face	100.0%	91.8%	97.4%
Iris	100.0%	90.8%	90.4%

**Table 2: Performance of multi-image attack: attack success rate and normal model accuracy on the Student model transferred from the infected Teacher and the clean Teacher.**

## 5.2 Results: Multi-Image Attack

Table 2 shows the attack performance on four tasks. We make two key observations. *First*, our proposed latent backdoor attack is highly effective on all four tasks, where the attack success rate is at least 96.6% if not 100%. This is particularly alarming since the attacker uses no more than 50 samples of the target ( $|X_{y_t}| \leq 50$ ) to infect the Teacher model, and can use generic images beyond  $X_{\setminus y_t}$  as adversarial inputs to the Student model.

*Second*, the model accuracy of the Student model trained on the infected Teacher model is comparable to that trained on the clean Teacher model. This means that the proposed latent backdoor attack does not compromise the model accuracy of the Student model (on clean inputs), thus the infected Teacher model is as attractive as its clean version.

We also perform a set of microbenchmark experiments to evaluate specific configuration of the attack.


**Microbenchmark 1: the need for trigger optimization.** As discussed in §4.3, a key element of our attack design is to compute the optimal trigger pattern  $\Delta_{opt}$  for  $y_t$ . We evaluate its effectiveness by comparing the attack performance of using randomly generated trigger patterns to that of using  $\Delta_{opt}$ .

Figure 5 shows the attack success rate vs. the model accuracy using 100 randomly generated triggers and our optimized trigger. Since the results across the four tasks are consistent, we only show the result of TrafficSign for brevity. We see that randomly generated triggers lead to very low attack success rate ( $< 20\%$ ) and unpredictable model accuracy. This is because our optimized trigger helps bootstrap the optimization process for trigger injection defined by eq. (4) to maximize the chance of convergence.

**Microbenchmark 2: the amount of non-target data  $X_{\setminus y_t}$ .** The key overhead of our proposed attack is to collect a set of target data  $X_{y_t}$ , and non-target data  $X_{\setminus y_t}$ , and use them to compute and inject the trigger into the Teacher model. In general  $|X_{\setminus y_t}| \gg |X_{y_t}|$ .

We experiment with different configurations of  $X_{\setminus y_t}$  by varying the number of classes and the number of instances per class. We arrive at two conclusions. *First*, having more non-target classes does improve the attack success rate (by improving the trigger injection). But the benefit of having more classes quickly converges, e.g., 8 out of 31 classes for Face and 32 out of 480 for Iris are sufficient to achieve 100% attack success rate. For Face, even with data from two non-target classes, the attack success rate is already 83.6%.

*Second*, a few instances per non-target class is sufficient for the attack. Again using Face as an example, 4 images per non-target class leads to 100% success rate while 2 images per class leads to



**Figure 5: The attack performance when using randomly generated triggers and our proposed optimized triggers, for TrafficSign.**

Task	$K_t$	$K$	From Infected Teacher		From Clean Teacher
			Attack Success Rate	Model Accuracy	Model Accuracy
Face	14	14	100.0%	91.8%	97.7%
	14	15	100.0%	91.4%	97.4%
	15	15	100.0%	94.0%	97.4%
Iris	14	14	100.0%	93.0%	94.4%
	14	15	100.0%	89.1%	90.4%
	15	15	100.0%	90.8%	90.4%

**Table 3: Performance of multi-image attack: attack success rate and normal model accuracy for different  $(K_t, K)$ .**

93.1% success rate. Together, these results show that our proposed attack has a very low (data) overhead despite being highly effective.

**Microbenchmark 3: the layer to inject the trigger.** As mentioned in §4.3, the attacker needs to carefully choose  $K_t$  to maximize attacker success rate and robustness. Our experiments show that for the given four tasks, the smallest  $K_t$  ( $K_t \leq K$ ) for a highly effective attack is either the first fully connected (FC) layer, e.g., 3 for Digit, 14 for Face and Iris, or the last convolutional layer, e.g., 6 for TrafficSign. Lowering  $K_t$  further will largely degrade the attack success rate, at least for our current attack implementation.

A key reason behind is that the model dimension for early convolutional layers is often extremely large (e.g., 25K for VGG-Face), thus the optimization defined by eq.(4) often fails to converge given the current data and computing resource. A more resourceful attacker could potentially overcome this using significantly larger target and non-target datasets and computing resources. We leave this to future work.

Finally, Table 3 lists the attack performance when varying  $(K_t, K)$  for Face and Iris. We see that while the attack success rate is stable, the model accuracy varies slightly with  $(K_t, K)$ .

## 5.3 Results: Single-image Attack

We now consider the extreme case where the attacker is only able to obtain a single image of the target, i.e.  $|X_{y_t}| = 1$ . For our evaluation, we reperform the above experiments but each time only use a single target image as  $X_{y_t}$ . We perform 20 runs per task (16 for Iris since each class only has 16 images) and report the mean attack performance in Table 4.



Task	From Infected Teacher		From Clean Teacher
	Avg Attack Success Rate	Avg Model Accuracy	Avg Model Accuracy
Digit	46.6%	97.5%	96.0%
TrafficSign	70.1%	83.6%	84.7%
Face	92.4%	90.2%	97.4%
Iris	78.6%	91.1%	90.4%

**Table 4: Performance of single-image attack.**

We make two key observations from these results. *First*, the attack success rate is lower than that of the multi-image attack. This is as expected since having only a single image of the target class makes it harder to accurately extract its features. *Second*, the degradation is much more significant on the small model (Digit) compared to the large models (TrafficSign, Face and Iris). We believe this is because larger models offer higher capacity (or freedom) to tune the feature representation by updating the model weights, thus the trigger can still be successfully injected into the Teacher model. In practice, the Teacher models designed for transfer learning are in fact large models, thus our proposed attack is highly effective with just a single image of the target.

## 6 REAL-WORLD ATTACK

So far, our experiments assume that the target data  $X_{y_t}$  for injecting latent backdoors comes from the same data source of the Student training data  $X_s$ . Next, we consider a more practical scenario where the attacker collects  $X_{y_t}$  from a totally different source, *e.g.*, by taking a picture of the physical target or searching for its images from the Internet.

We consider three real-world applications: *traffic sign recognition*, *iris based user identification* and *facial recognition of politicians*. We show that the attacker can successfully launch latent backdoor attacks against these applications and cause harmful misclassification events, by just using pictures taken by commodity smartphones or those found from Google Image and Youtube. Again, our experiments assume that  $K_t = K$ .

### 6.1 Ethics and Data Privacy

Our experiments are designed to reproduce the exact steps a real-world attack would entail. However, we are very aware of the sensitive nature of some of these datasets. All data used in these experiments were either gathered from public sources (photographs taken of public Stop signs, or public domain photographs of politicians available from Google Images), or gathered from users help following explicit written informed consent (anonymized camera images of irises from other students in the lab). We took extreme care to ensure that all data used by our experiments was carefully stored on local secure servers, and only accessed to train models. Our iris data will be deleted once our experimental results are finalized.

### 6.2 Traffic Sign Recognition

Real-world attack on traffic sign recognition, if successful, can be extremely harmful and create life-threatening accidents. For example, the attacker can place a small sticker (*i.e.* the trigger) on a speed limit sign, causing nearly self-driving cars to misclassify it into a stop sign and stop abruptly in the middle of the road. To

Scenario	Multi-image Attack		Single-image Attack	
	Attack Success Rate	Model Accuracy	Avg Attack Success Rate	Avg Model Accuracy
Traffic Sign	100%	88.8%	67.1%	87.4%
Iris Identification	90.8%	96.2%	77.1%	97.7%
Politician Recognition	99.8%	97.1%	90.0%	96.7%

**Table 5: Attack performance in real-world scenarios.**

launch a conventional backdoor attack against this application (*e.g.*, via BadNets [19]), the attacker needs to have access to the self-driving car’s model training data and/or control its model training.

Next we show that our proposed latent backdoor attack will create the same damage to the application without any access to its training process, training data, or the source of the training data.

**Attack Configuration.** The attacker uses the public available Germany traffic sign dataset (*e.g.*, GTSRB) to build the (clean) Teacher model. To inject the latent backdoor trigger, the attacker uses a subset of the GTSRB classes as the non-target data ( $X_{\setminus y_t}$ ). To form the target data  $X_{y_t}$  (*i.e.* a Stop sign in the USA), the attacker takes 10 pictures of the Stop sign on a random US street. Figure 6 shows a few examples we took with commodity smartphones. The attacker then releases the Teacher model and waits for any victim to download the model and use transfer learning to build an application on US traffic sign recognition.

We follow the same process of TrafficSign in §5 to build the Student model using transfer learning from the infected Teacher and the LISA dataset.

**Attack Performance.** Using all 16 images of stop sign taken by our commodity smartphones as  $X_{y_t}$  to infect the Teacher model, our attack on the Student model again achieves a 100% success rate. Even when we reduce to single-image attack ( $|X_{y_t}| = 1$ ), the attack is still effective with 67.1% average success rate (see Table 5).

### 6.3 Iris Identification

The attacker seeks to get physical access to a company’s building that will use iris recognition for user identification in the near future. The attacker also knows that the target  $y_t$  will be a legitimate user in this planned iris recognition system. Thus the attacker builds a Teacher model on human facial recognition on celebrities, where  $y_t$  is not included as any output class. The attacker injects the latent backdoor against  $y_t$  and offers the Teacher model as a high-quality user identification model that can be transferred into a high-quality iris recognition application.

**Attack Configuration.** Like Face, the attacker starts from the VGG-Face model as a clean Teacher model, and forms the non-target data  $X_{\setminus y_t}$  using the CASIA IRIS dataset, which is publicly available. To build the target data  $X_{y_t}$ , the attacker searches for  $y_t$ ’s headshots on Google, and crops out the iris area of the photos. The final  $X_{y_t}$  consists of 5 images of the target  $y_t$  (images omitted for privacy protection).

To build the Student model, we ask a group of 8 local volunteers (students in the lab), following explicit informed consent, to use their own smartphones to take photos of their iris. The resulting training data  $X_s$  used by transfer learning includes 160 images from



Figure 6: Pictures of real-world stop signs as  $X_{y_t}$  which we took using a smartphone camera. Part of the image is modified for anonymization purpose.

8 people. In this case,  $X_{y_t}$ ,  $X_{\setminus y_t}$  and  $X_s$  all come from different sources.

**Attack Performance.** Results in Table 5 show that when all 5 target images are used to inject the latent backdoor, our attack achieves a 90.8% success rate. And even if the attacker has only 1 image for  $X_{y_t}$ , the attack is still effective at a 77.1% success rate.

#### 6.4 Facial Recognition on Politicians

Finally, we demonstrate a scenario where the attacker leverages the chronological advantage of the latent backdoor attack. Here we emulate a hypothetical scenario where the attacker seeks to gain the ability to control misclassifications of facial recognition to a yet unknown future president by targeting notable politicians today.

Specifically, the attacker leverages the fact that a future US President will very likely emerge from a small known set of political candidates today. The attacker builds a high-quality Teacher model on face recognition of celebrities, and injects a set of latent backdoors targeting presidential candidates. The attacker actively promotes the Teacher model for adoption (or perhaps leverages an insider to alter the version of the Teacher model online). A few months later, a new president is elected (out of one of our likely presidential candidates), the White House team adds the president’s facial images into its facial recognition system, using a Student model derived from our infected Teacher model. This activates our latent backdoor, turning it into a live backdoor attack. As the facial recognition system is built prior to the current presidential election, it is hard for the White House team to think about the possibility of any backdoors, and any checks on the Teacher model will not reveal any unexpected or unusual behavior.

**Attack Configuration.** Similar to the Face task in §5, the attacker uses the VGG-Face model as the clean Teacher model and the VGG-Face dataset as the non-target dataset  $X_{\setminus y_t}$ . The attacker selects 9 top leaders as targets and collects their (low-resolution) headshots from Google. The resulting  $X_{y_t}$  will include 10 images per target for 9 targets, and a total of 90 images. Some examples for a single target are shown in Figure 7.

To train the Student model, we assume the White House team uses its own source rather than VGG-Face. We emulate this using a set of high-resolution videos of Congress members from Youtube,



Figure 7: Examples of target politician images that we collected as  $X_{y_t}$ .

from which we extract multiple headshot frames from each person’s video. The resulting dataset is 1.7K images in 13 classes.

**Performance of Single- and Multi-target Attacks.** Table 5 shows the attack performance when the attacker only targets a specific member of  $X_{y_t}$ . The success rate is 99.8% for multi-image attack (using all 10 images) and 90.0% for single-image attack (averaged over the 10 images).

Since it is hard to guess the future president, the attacker increases its attack success rate by injecting latent backdoors of multiple targets into the Teacher model. Figure 8 plots the attack performance as we vary the number of targets. We see that the attack success rate stays close to 100% when injecting up to 3 targets, and then drops gracefully as we add more targets. But even with 9 targets, the success rate is still 60%. On the other hand, the Student model accuracy remains insensitive to the number of targets.

The trend that the attack success rate drops with the number of targets is as expected, and the same trend is observed on conventional backdoor attacks [50]. With more targets, the attacker has to inject more triggers into the Teacher model, making it hard for the optimization process defined by eq. (4) to reach convergence. Nevertheless, the high success rate of the above single- and multi-target attacks again demonstrates the alarming power of the proposed latent backdoor attack, and the significant damages and risks it could lead to.

## 7 DEFENSE

In this section, we explore and evaluate potential defenses against our attack. Our discussion below focuses on the Face task described in §5.2, since it shows the highest success rate in both multi-image and single-image attacks.

### 7.1 Leveraging Existing Backdoor Defenses

Our first option is to leverage existing defenses proposed for normal backdoor attacks. We consider two state-of-the-art defenses: Neural Cleanse [50] and Fine-Pruning [29] (as discussed in §2.1). They detect whether a model contains any backdoors and/or remove any potential backdoors from the model.




Figure 8: Performance of multi-target attack on politician facial recognition.

**Neural Cleanse.** Neural Cleanse [50] cannot be applied to a Teacher model, because it requires access to the label of the target ( $y_t$ ). Instead, we run it on an infected Student model along with the Student training data. When facing conventional backdoor attacks (e.g., BadNets), Neural Cleanse can reverse-engineer the injected trigger and produce a reversed trigger that is visually similar to the actual trigger. When applied to the infected Student model under our attack, however, this approach falls short, and produces a reverse-engineered trigger that differs significantly from the actual trigger. Our intuition says that Neural Cleanse fails under because trigger reverse-engineering is based on end-to-end optimization from the input space to the final label space. It is unable to detect any manipulation that terminates at an intermediate feature space.

**Fine-Pruning.** Fine-Pruning [29] can be used to disrupt potential backdoor attacks, but is “blind,” in that it does not detect whether a model has a backdoor installed. Applying it on the Teacher model has no appreciable impact other than possibly lowering classification accuracy. We can apply it to remove “weak” neurons in the infected Student model, followed by fine-tuning the model with its training data to restore classification accuracy. Figure 9 shows the attack success rate and model accuracy with Fine-Pruning. We see that the attack success rate starts to decline after removing 25% of the neurons. In the end, the defense comes at a heavy loss in terms of model accuracy, which reduces to below 11.5%. Thus Fine-Pruning is not a practical defense against our latent backdoors.

## 7.2 Input Image Blurring

As mentioned in §5.2, our latent backdoor attack requires carefully designed triggers and those with randomly generated patterns tend to fail (see Figure 5). Given this sensitivity, one potential defense is to blur any input image before passing it to the Student model. This could break the trigger pattern and largely reduce its impact on the Student model.

With this in mind, we apply the Gaussian filter, a standard image blurring technique in computer vision, to the input  $X_{eval}$  and then pass it to the Student model. Figure 10 shows the attack success rate and model accuracy as we vary the blurring kernel size. The larger the kernel size is, the more blurred the input image becomes. Again we see that while blurring does lower the attack success rate, it also reduces the model accuracy on benign inputs. Unlike




Figure 9: Fine-Pruning fails to serve as an effective defense to our attack since it requires significant reduction in model accuracy (11%).

Fine-Pruning, here the attack success rate drops faster than the model accuracy. Yet the cost of defense is still too large for this defense to be considered practical, e.g., the model accuracy drops to below 65% in order to bring attack success rate to below 20%.

## 7.3 Multi-layer Tuning in Transfer Learning

The final defense leverages the fact that the attacker is unable to control the exact set of layers that the transfer learning will update. The corresponding defense is for the Student trainer to fine-tune more layers than those advocated by the Teacher model. Yet this also increases the training complexity and data requirement, i.e. more training data is required for the model to converge.

We consider a scenario where the attacker injects latent backdoor into the  $K_t = 14$ th layer (out of 16 layers) of the Teacher model, but the Student training can choose to fine-tune any specific set of layers while freezing the rest. Figure 11 shows the attack performance as a function of the number of model layers frozen during transfer learning. 0 means no layers are frozen, i.e. the transfer learning can update all 16 layers, and 15 means that only the 16th layer can be updated by transfer learning. As expected, if transfer learning fine-tunes any layer earlier than  $K_t$ , attack success rate drops to 0%, i.e. the trigger gets wiped out.

It should be noted that since the Student has no knowledge of  $K_t$ , the ideal defense is to fine-tune all layers in the Teacher model. Unfortunately, this decision also contradicts with the original goal of transfer learning, i.e. using limited training data to build an accurate model. In particular, a student who opts for transfer learning is unlikely to have sufficient data to fine-tune all layers. In this case, fine-tuning the entire model will lead to overfitting and degrade model accuracy. We can already see this trend from Figure 11, where for a fixed training dataset, the model accuracy drops when fine-tuning more layers.

Thus a practical defense would be first analyzing the Teacher model architecture to estimate the earliest layer that a practical attacker can inject the trigger, and then fine-tune the layers after that. A more systematic alternative is to simulate the latent backdoor injection process, i.e. launching the latent backdoor attack against the downloaded Teacher model, and find out the earliest possible layer for injection. However, against a powerful attacker capable of injecting the latent backdoor at an earlier layer, the defense would




Figure 10: Input blurring is not a practical defense since it still requires heavy drop of model accuracy to reduce attack success rate.

need to incur the cost of fine-tuning more layers, potentially all layers in the model.

## 8 RELATED WORK

**Other Backdoor Attacks and Defenses.** In addition to attacks mentioned in §2.1, Chen *et al.* proposed a backdoor attack under a more restricted scenario, where the attacker can only pollute a limited portion of training set [13]. Another line of work directly tampers with the hardware a DNN model runs on [15, 28]. Such backdoor circuits could also affect the model performance when a trigger is present. Our proposed attack differs by not requiring any access to the Student model, its data or operating hardware.

In terms of defenses, Liu *et al.* [31] only presented some brief intuitions on backdoor detection, while Chen *et al.* [13] reported a number of ideas that are shown to be ineffective. Liu *et al.* [32] proposed three defenses: input anomaly detection, re-training, and input preprocessing, and require the poisoned training data. A more recent work [49] leveraged trace in the spectrum of the covariance of a feature representation to detect backdoor. It also requires the poisoned training data. Like Neural Cleanse and Fine-Pruning, these defenses only target normal backdoor attack and cannot be applied to our latent backdoor attack.

**Transfer Learning.** In a deep learning context, transfer learning has been shown to be effective in vision [6, 12, 44, 45], speech [14, 21, 25, 52], and text [23, 35]. Yosinski *et al.* compared different transfer learning approaches and studied their impact on model performance [54]. Razavian *et al.* studied the similarity between Teacher and Student tasks, and analyzed its correlation with model performance [43].

**Poisoning Attacks.** Poisoning attack pollutes training data to alter a model’s behavior. Different from backdoor attack, it does not rely on any trigger, and manipulates the model’s behavior on a set of clean samples. Defenses against poisoning attack mostly focus on sanitizing the training set and removing poisoned samples [7, 16, 22, 38, 46, 48]. The insight is to find samples that would alter model’s performance significantly [7]. This insight is less effective against backdoor attack [13], as injected samples do not affect the model’s performance on clean samples. It is also impractical under our




Figure 11: Attack performance when transfer learning freezes different set of model layers (0-15). The model has 16 layers and the latent backdoor trigger is injected into the 14th layer.

attack model, as the defender does not have access to the poisoned training set (used by the Teacher).

**Other Adversarial Attacks against DNNs.** Numerous (non-backdoor) adversarial attacks have been proposed against general DNNs, often crafting imperceptible modifications to images to cause misclassification. These can be applied to DNNs during inference [10, 26, 30, 39, 51]. A number of defenses have been proposed [24, 33, 34, 40, 53], yet many have shown to be less effective against an adaptive adversary [4, 8, 9, 20]. Recent works tried to craft universal perturbations that trigger misclassification for multiple images in an uninfected DNN [5, 37].

## 9 CONCLUSION

In this paper, we identify a new, more powerful variant of the backdoor attack against deep neural networks. Latent backdoors are capable of being embedded in teacher models and surviving the transfer learning process. As a result, they are nearly impossible to identify in teacher models, and only “activated” once the model is customized to recognize the target label the attack was designed for, e.g. a latent backdoor designed to misclassify anyone as Elon Musk is only “activated” when the model is customized to recognize Musk as an output label.

We demonstrate the effectiveness and practicality of latent backdoors through extensive experiments and real-world tests. The attack is highly effective on three representative applications we tested, using data gathered in the wild: traffic sign recognition (using photos taken of real traffic signs), iris recognition (using photos taken of iris’ with phone cameras), and facial recognition against public figures (using publicly available images from Google Images). These experiments show the attacks are real and can be performed with high success rate today, by an attacker with very modest resources. Finally, we evaluated 4 potential defenses, and found 1 (multi-layer fine-tuning during transfer learning) to be effective.

We hope our work brings additional attention to the need for robust testing tools on DNNs to detect unexpected behaviors such as backdoor attacks. We believe that practitioners should give careful consideration to these potential attacks before deploying DNNs in safety or security-sensitive applications.



## REFERENCES

- [1] 2010. <http://biometrics.idealtest.org/>. (2010). CASIA Iris Dataset.
- [2] 2017. [http://pytorch.org/tutorials/beginner/transfer\\_learning\\_tutorial.html](http://pytorch.org/tutorials/beginner/transfer_learning_tutorial.html). (2017). PyTorch transfer learning tutorial.
- [3] 2017. <https://codalabs.developers.google.com/codalabs/cpb102-txf-learning/index.html>. (2017). Image Classification Transfer Learning with Inception v3.
- [4] Anish Athalye, Nicholas Carlini, and David Wagner. 2018. Obfuscated gradients give a false sense of security: Circumventing defenses to adversarial examples. In *Proc. of ICML*.
- [5] Tom B Brown, Dandelion Mané, Aurko Roy, Martin Abadi, and Justin Gilmer. 2017. Adversarial patch. *arXiv preprint arXiv:1712.09665* (2017).
- [6] Sergi Caelles, Kevis-Kokitsi Maninis, Jordi Pont-Tuset, Laura Leal-Taixé, Daniel Cremers, and Luc Van Gool. 2017. One-shot video object segmentation. In *Proc. of CVPR*.
- [7] Yinzhi Cao, Alexander Fangxiao Yu, Andrew Aday, Eric Stahl, Jon Merwine, and Junfeng Yang. 2018. Efficient Repair of Polluted Machine Learning Systems via Causal Unlearning. In *Proc. of ASIACCS*.
- [8] Nicholas Carlini and David Wagner. 2016. Defensive distillation is not robust to adversarial examples. *arXiv preprint arXiv:1607.04311* (2016).
- [9] Nicholas Carlini and David Wagner. 2017. Magnet and efficient defenses against adversarial attacks are not robust to adversarial examples. *arXiv preprint arXiv:1711.08478* (2017).
- [10] Nicholas Carlini and David Wagner. 2017. Towards evaluating the robustness of neural networks. In *Proc. of IEEE S&P*.
- [11] Bryant Chen, Wilka Carvalho, Nathalie Baracaldo, Heiko Ludwig, Benjamin Edwards, Taesung Lee, Ian Molloy, and Biplav Srivastava. 2018. Detecting backdoor attacks on deep neural networks by activation clustering. *arXiv preprint arXiv:1811.03728* (2018).
- [12] Jun-Cheng Chen, Rajeev Ranjan, Amit Kumar, Ching-Hui Chen, Vishal M Patel, and Rama Chellappa. 2015. An end-to-end system for unconstrained face verification with deep convolutional neural networks. In *Proc. of Workshop on ICCV*.
- [13] Xinyun Chen, Chang Liu, Bo Li, Kimberly Lu, and Dawn Song. 2017. Targeted Backdoor Attacks on Deep Learning Systems Using Data Poisoning. *arXiv preprint arXiv:1712.05526* (2017).
- [14] Dan C Cireşan, Ueli Meier, and Jürgen Schmidhuber. 2012. Transfer learning for Latin and Chinese characters with deep neural networks. In *Proc of IJCNN*.
- [15] Joseph Clements and Yingjie Lao. 2018. Hardware Trojan Attacks on Neural Networks. *arXiv preprint arXiv:1806.05768* (2018).
- [16] Gabriela F Cretu, Angelos Stavrou, Michael E Locasto, Salvatore J Stolfo, and Angelos D Keromytis. 2008. Casting out demons: Sanitizing training data for anomaly sensors. In *Proc. of IEEE S&P*.
- [17] Kevin Eykholt, Ivan Evtimov, Earlene Fernandes, Bo Li, Amir Rahmati, Chaowei Xiao, Atul Prakash, Tadayoshi Kohno, and Dawn Song. 2018. Robust physical-world attacks on deep learning models. In *Proc. of CVPR*.
- [18] Tianyu Gu, Brendan Dolan-Gavitt, and Siddharth Garg. 2017. Badnets: Identifying vulnerabilities in the machine learning model supply chain. In *Proc. of Machine Learning and Computer Security Workshop*.
- [19] Tianyu Gu, Kang Liu, Brendan Dolan-Gavitt, and Siddharth Garg. 2019. BadNets: Evaluating Backdooring Attacks on Deep Neural Networks. *IEEE Access* (2019).
- [20] Warren He, James Wei, Xinyun Chen, Nicholas Carlini, and Dawn Song. 2017. Adversarial example defenses: Ensembles of weak defenses are not strong. In *Proc. of WOOT*.
- [21] Georg Heigold, Vincent Vanhoucke, Alan Senior, Patrick Nguyen, Marc’Aurélio Ranzato, Matthieu Devin, and Jeffrey Dean. 2013. Multilingual acoustic models using distributed deep neural networks. In *Proc. of ICASSP*.
- [22] Matthew Jagielski, Alina Oprea, Battista Biggio, Chang Liu, Cristina Nita-Rotaru, and Bo Li. 2018. Manipulating machine learning: Poisoning attacks and counter-measures for regression learning. In *Proc. of IEEE S&P*.
- [23] Melvin Johnson, Mike Schuster, Quoc V Le, Maxim Krikun, Yonghui Wu, Zhifeng Chen, Nikhil Thorat, Fernanda Viégas, Martin Wattenberg, Greg Corrado, and others. 2017. Google’s multilingual neural machine translation system: enabling zero-shot translation. In *Proc. of ACL*.
- [24] Harini Kannan, Alexey Kurakin, and Ian Goodfellow. 2018. Adversarial Logit Pairing. *arXiv preprint arXiv:1803.06373* (2018).
- [25] Julius Kunze, Louis Kirsch, Ilia Kurenkov, Andreas Krug, Jens Johannsmeier, and Sebastian Stober. 2017. Transfer learning for speech recognition on a budget. In *Proc. of RePL4NLP*.
- [26] Alexey Kurakin, Ian Goodfellow, and Samy Bengio. 2017. Adversarial machine learning at scale. In *Proc. of ICLR*.
- [27] Yann LeCun, LD Jackel, Léon Bottou, Corinna Cortes, John S Denker, Harris Drucker, Isabelle Guyon, Urs A Muller, Eduard Sackinger, Patrice Simard, and others. 1995. Learning algorithms for classification: A comparison on hand-written digit recognition. *Neural networks: the statistical mechanics perspective* (1995).
- [28] Wenshuo Li, Jincheng Yu, Xuefei Ning, Pengjun Wang, Qi Wei, Yu Wang, and Huazhong Yang. 2018. Hu-Fu: Hardware and Software Collaborative Attack Framework against Neural Networks. In *Proc. of ISVLSI*.
- [29] Kang Liu, Brendan Dolan-Gavitt, and Siddharth Garg. 2018. Fine-Pruning: Defending Against Backdooring Attacks on Deep Neural Networks. In *Proc. of RAID*.
- [30] Yanpei Liu, Xinyun Chen, Chang Liu, and Dawn Song. 2016. Delving into transferable adversarial examples and black-box attacks. In *Proc. of ICLR*.
- [31] Yingqi Liu, Shiqing Ma, Yousra Aafer, Wen-Chuan Lee, Juan Zhai, Weihang Wang, and Xiangyu Zhang. 2018. Trojaning attack on neural networks. In *Proc. of NDSS*.
- [32] Yuntao Liu, Yang Xie, and Ankur Srivastava. 2017. Neural trojans. In *Proc. of ICCD*.
- [33] Aleksander Madry, Aleksandar Makelov, Ludwig Schmidt, Dimitris Tsipras, and Adrian Vladu. 2018. Towards deep learning models resistant to adversarial attacks. In *Proc. of ICLR*.
- [34] Dongyu Meng and Hao Chen. 2017. Magnet: a two-pronged defense against adversarial examples. In *Proc. of CCS*.
- [35] Tomas Mikolov, Quoc V Le, and Ilya Sutskever. 2013. Exploiting similarities among languages for machine translation. *arXiv preprint arXiv:1309.4168* (2013).
- [36] Andreas Mogelmose, Mohan Manubhai Trivedi, and Thomas B Moeslund. 2012. Vision-based traffic sign detection and analysis for intelligent driver assistance systems: Perspectives and survey. *IEEE Transactions on Intelligent Transportation Systems* 13, 4 (2012).
- [37] Seyed-Mohsen Moosavi-Dezfooli, Alhussein Fawzi, Omar Fawzi, and Pascal Frossard. 2017. Universal adversarial perturbations. In *Proc. of CVPR*.
- [38] Mehran Mozaffari-Kermani, Susmita Sur-Kolay, Anand Raghunathan, and Niraj K Jha. 2015. Systematic poisoning attacks on and defenses for machine learning in healthcare. *IEEE journal of biomedical and health informatics* 19, 6 (2015), 1893–1905.
- [39] Nicolas Papernot, Patrick McDaniel, Somesh Jha, Matt Fredrikson, Z Berkay Celik, and Ananthram Swami. 2016. The limitations of deep learning in adversarial settings. In *Proc. of Euro S&P*.
- [40] Nicolas Papernot, Patrick McDaniel, Xi Wu, Somesh Jha, and Ananthram Swami. 2016. Distillation as a defense to adversarial perturbations against deep neural networks. In *Proc. of IEEE S&P*.
- [41] Omkar M Parkhi, Andrea Vedaldi, and Andrew Zisserman. 2015. Deep face recognition. In *Proc. of BMVC*.
- [42] Nicolas Pinto, Zak Stone, Todd Zickler, and David Cox. 2011. Scaling up biologically-inspired computer vision: A case study in unconstrained face recognition on facebook. In *Proc. of CVPR Workshop*.
- [43] Ali Sharif Mozavian, Hossein Azizpour, Josephine Sullivan, and Stefan Carlsson. 2014. CNN features off-the-shelf: an astounding baseline for recognition. In *Proc. of Workshop on CVPR*.
- [44] Joseph Redmon, Santosh Divvala, Ross Girshick, and Ali Farhadi. 2016. You only look once: Unified, real-time object detection. In *Proc. of CVPR*.
- [45] Shaoqing Ren, Kaiming He, Ross Girshick, and Jian Sun. 2015. Faster R-CNN: Towards real-time object detection with region proposal networks. In *Proc. of NeurIPS*.
- [46] Benjamin IP Rubinstein, Blaine Nelson, Ling Huang, Anthony D Joseph, Shing-hon Lau, Satish Rao, Nina Taft, and JD Tygar. 2009. Antidote: understanding and defending against poisoning of anomaly detectors. In *Proc. of IMC*.
- [47] Johannes Stalkamp, Marc Schlipf, Jan Salmen, and Christian Igel. 2011. The German Traffic Sign Recognition Benchmark: A multi-class classification competition. In *Proc. of IJCNN*.
- [48] Jacob Steinhardt, Pang Wei W Koh, and Percy S Liang. 2017. Certified defenses for data poisoning attacks. In *Proc. of NeurIPS*.
- [49] Brandon Tran, Jerry Li, and Aleksander Madry. 2018. Spectral signatures in backdoor attacks. In *Proc. of NeurIPS*.
- [50] Bolun Wang, Yuanshun Yao, Shawn Shan, Huiying Li, Bimal Viswanath, Haitao Zheng, and Ben Y Zhao. 2019. Neural cleanse: Identifying and mitigating backdoor attacks in neural networks. In *Proc. of IEEE S&P*.
- [51] Bolun Wang, Yuanshun Yao, Bimal Viswanath, Zheng Haitao, and Ben Y. Zhao. 2018. With Great Training Comes Great Vulnerability: Practical Attacks against Transfer Learning. In *Proc. of USENIX Security*.
- [52] Dong Wang and Thomas Fang Zheng. 2015. Transfer learning for speech and language processing. In *Proc. of APSIPA*.
- [53] Weilin Xu, David Evans, and Yanjun Qi. 2018. Feature squeezing: Detecting adversarial examples in deep neural networks. In *Proc. of NDSS*.
- [54] Jason Yosinski, Jeff Clune, Yoshua Bengio, and Hod Lipson. 2014. How transferable are features in deep neural networks?. In *Proc. of NeurIPS*.

## APPENDIX

**Model Architecture.** Table 6, 7, and 8 list the detailed architecture of the Teacher model for the four applications considered by our evaluation in §5. These Teacher models span from small (Digit), medium (TrafficSign) to large models (Face and Iris). We also list the index of every layer in each model. Note that the index of pooling layer is counted as its previous layer, as defined conventionally.

**Table 6: Teacher model architecture for Digit. FC stands for fully-connected layer. Pooling layer’s index is counted as its previous layer.**

Layer Index	Layer Type	# of Channels	Filter Size	Stride	Activation
1	Conv	16	5×5	1	ReLU
1	MaxPool	16	2×2	2	-
2	Conv	32	5×5	1	ReLU
2	MaxPool	32	2×2	2	-
3	FC	512	-	-	ReLU
4	FC	5	-	-	Softmax

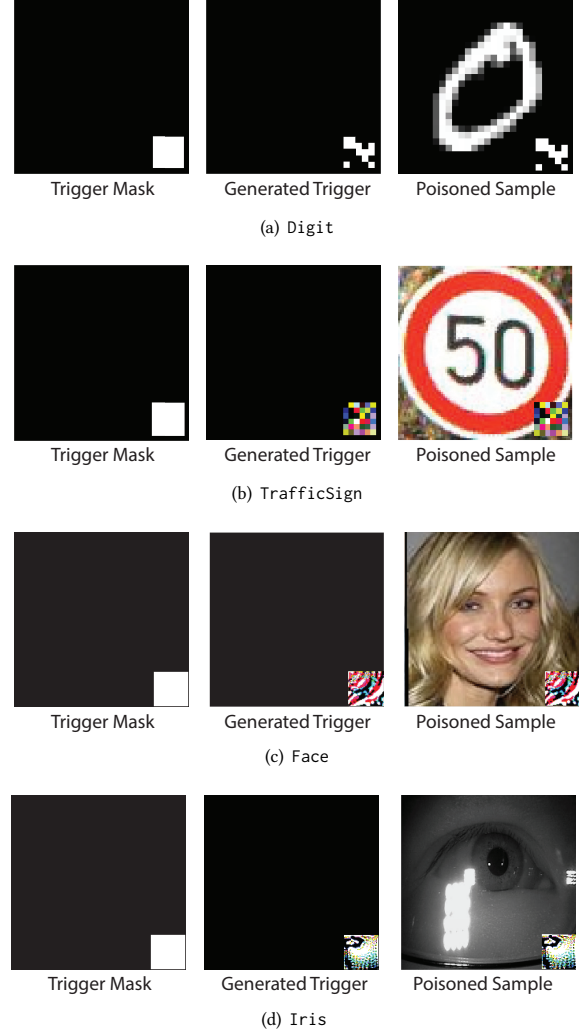
**Table 7: Teacher model architecture for TrafficSign.**

Layer Index	Layer Type	# of Channels	Filter Size	Stride	Activation
1	Conv	32	3×3	1	ReLU
2	Conv	32	3×3	1	ReLU
2	MaxPool	32	2×2	2	-
3	Conv	64	3×3	1	ReLU
4	Conv	64	3×3	1	ReLU
4	MaxPool	64	2×2	2	-
5	Conv	128	3×3	1	ReLU
6	Conv	128	3×3	1	ReLU
6	MaxPool	128	2×2	2	-
7	FC	512	-	-	ReLU
8	FC	43	-	-	Softmax

**Table 8: Teacher model architecture for Face and Iris.**

Layer Index	Layer Type	# of Channels	Filter Size	Stride	Activation
1	Conv	64	3×3	1	ReLU
2	Conv	64	3×3	1	ReLU
2	MaxPool	64	2×2	2	-
3	Conv	128	3×3	1	ReLU
4	Conv	128	3×3	1	ReLU
4	MaxPool	128	2×2	2	-
5	Conv	256	3×3	1	ReLU
6	Conv	256	3×3	1	ReLU
7	Conv	256	3×3	1	ReLU
7	MaxPool	256	2×2	2	-
8	Conv	512	3×3	1	ReLU
9	Conv	512	3×3	1	ReLU
10	Conv	512	3×3	1	ReLU
10	MaxPool	512	2×2	2	-
11	Conv	512	3×3	1	ReLU
12	Conv	512	3×3	1	ReLU
13	Conv	512	3×3	1	ReLU
13	MaxPool	512	2×2	2	-
14	FC	4096	-	-	ReLU
15	FC	4096	-	-	ReLU
16	FC	2622	-	-	Softmax

**Target-dependent Trigger Generation.** Figure 12 shows samples of backdoor triggers generated by our attacks as discussed in §5. The trigger mask is chosen to be a square-shaped pattern located at the bottom right of each input image. The trigger generation process maximizes the trigger effectiveness against  $y_t$  by minimizing the difference between poisoned non-target samples and clean target samples described by eq. (2). These generated triggers are used to inject latent backdoor into the Teacher model. They are also used to launch misclassification attacks after any Student model is trained from the infected Teacher model.



**Figure 12: Samples of triggers produced by our attack and the corresponding poisoned images.**